

TEAM PREP · FINAL PRESENTATION

RetainIQ — Autonomous Workforce Retention Engine

We are judged on the **prototype** — not on slides, not on extra features.

🕒 5 MINUTES TOTAL · then judges ask anything

🎯 4 criteria · 1 working product

1 What the judges are scoring

Criterion	What that means here	Where we prove it
1) Innovation	Is the idea new and does the prototype show it, not just claim it?	The closed loop: diagnosis → auto-executed, owned, dated plan
2) Technically feasible	Does it actually run, with real data and real security?	Live agent call, real database, RLS, working auth, automated tests
3) Presentation skills	Clear 5-minute story, confident demo, no fluff.	This run-of-show; everyone knows their role
4) Critical thinking	Do we know our limits, and can we defend trade-offs?	Honest "not connected" boundaries + next steps + Q&A bank

2 One-line pitch (say this in the first 15 seconds)

“RetainIQ reads each employee’s real signals, uses an AI agent to diagnose flight risk, and automatically turns that diagnosis into an owned, dated retention plan that the team executes with evidence.”

Judges want the prototype to *do* this, live. So the 5 minutes are ~80% demo.

3 The 5-minute run-of-show (memorize these time stamps)

Time	Beat	What happens on screen
0:00-0:30	Hook + pitch	One-line pitch. No slides. Point at the product.
0:30-1:30	Problem → the loop	“Flight risk is diagnosed but nothing happens. RetainIQ closes the loop: <i>diagnose</i> → <i>plan</i> → <i>execute</i> .”
1:30-4:00	LIVE DEMO (the meat)	Sign in → Dashboard → Deep-Dive on the highest-risk employee → run the agent → Execute retention plan → Action Center (tasks, owners, due dates).
4:00-4:40	Innovation + honesty	2 sentences on why this is different. Then: “What is <i>not</i> connected and why.”
4:40-5:00	Close	“This is a real, working loop — we’re happy to take questions.”

Hard rule: the demo is the presentation. If something is still on-screen when the timer hits 5:00, finish the sentence and invite questions. Never apologize for time; never “I’ll show more if we have time.”

4 Demo script — click for click

ONE DEMOER ON THE MACHINE, ONE NARRATOR. PRACTICE THIS EXACT ORDER UNTIL IT'S MUSCLE MEMORY.

Step	Action	Say (talking point)
1 • Sign in	Log in with the prepared account (or show the alternate auth mode).	"Real authentication — every screen is role-scoped."
2 • Dashboard	Show live stats and the highest-risk employees.	"Live signals from the real database, not a mock."
3 • Deep-Dive	Open the top-risk employee. Point at attendance/workload, sentiment, skills.	"One employee, one record, real telemetry."
4 • Run the agent	Click Run Qwen diagnostic . Narrate while it works (~60s). Show the live "agent activity" panel.	"A published AI agent reviews only this employee's stored record. We show its reasoning and tool activity — no black box."
5 • Read the result	Walk through summary → observations → hypotheses → recommended actions (each with an owner).	"Structured, schema-validated output — the app rejects anything that doesn't match, so we never show a made-up result."
6 • Execute	Click Execute retention plan → it jumps to the Action Center.	"One click turns the diagnosis into a tracked case — the agent's actions become tasks with owners and staggered due dates."
7 • Action Center	Show the plan summary, the "AUTO-ORCHESTRATED PLAN" label, task owners, due dates.	"This is the automatic execution — real, stored, accountable. No fake orchestration."
8 • (Optional, time)	Record one completion with an evidence note; open its history log.	"Completing a task requires human evidence. The history is append-only."

If the agent is slow or fails (this WILL happen at some point — stay calm)

While it runs:

- Narrate what's happening instead of staring at the spinner: "the agent has a 180-second cap and one automatic retry."
- Show the reasoning/tool-activity panel — it's a feature, not a wait screen.

If it errors:

- Say it honestly: "The live model hit its limit — the app refuses to fake a result. Here is the validated output from a real run we captured earlier."
- Open the **copied JSON** from a real run (the "Copy JSON" button) and continue the flow.
- Never claim the failure is part of the demo. Honesty is a scoring point (critical thinking).

5 Proving the 4 criteria with the prototype

1) Innovation — “the loop is closed”

- Most tools *predict* attrition. RetainIQ *executes*: the diagnosis **becomes** an owned, dated plan (case + tasks + owners + staggered due dates).
- Trustworthy AI: strict output schema, alias normalization, balanced-JSON extraction, auto-repair, employee-mismatch rejection, one auto-retry, transparent agent activity.
- Honest-by-design: execution is real inside the app; external dispatch is labeled “not connected” instead of being faked.

2) Technically feasible — “it runs, for real”

- **Real backend**: database, authentication, and backend functions; Row-Level Security on every table; all mutations go through server-side functions (direct client writes are denied).
- **Real AI call** end-to-end to a published agent — validated live in tests.
- **Real security**: privilege-escalation attempts are blocked (forged thread ownership, direct task writes).
- **Real quality**: automated browser suite (6 tests) covering auth quality, AA contrast in both themes, mobile performance (auth load $\approx 2.3s$), the full authenticated journey, and RLS denials.

3) Presentation skills — “we show, we don’t tell”

- Roles: **Narrator** (the story), **Demoer** (the clicks), **Timekeeper** (hard stop at 5:00), **Shadow** (agent backup + Q&A support).
- Every claim is backed by a click we can make in the room.
- Short sentences, present tense, one idea per screen.

4) Critical thinking — “we know our limits”

- We state what is **not** connected (manager dispatch, training enrollment, HRIS updates) and why we chose truthfulness over fake integrations.
- The agent is **decision support**, not an employment decision-maker — humans confirm with evidence.
- We acknowledge real failure modes: agent latency/variance, timeouts, per-user case rules, RLS boundaries.
- We can say what we’d build next and why it matters (see next section).

6 If they ask “what’s next?” — have these ready

Next step	Why it matters (in one line)
Connect HRIS / LMS / email dispatch via backend functions	Turns “execute” into real external actions with the same evidence trail.
Due-date reminders / notifications	Closes the loop over time, not just at creation.
Org-level attrition analytics	From one employee to workforce patterns.
Model swap + fine-tuning per company	Better diagnoses per industry via the same gateway.
Explainability/export of the full case history	Audit-ready retention decisions.

7 Q&A bank — answer any question in ≤3 sentences

ANSWER WITH THE PROTOTYPE FIRST. “IN THE PRODUCT, ...” · “WE CHOSE ... BECAUSE ...” · ONE OF THESE IS ALMOST CERTAINLY ASKED.

Q: Is this a real product or a mockup?

Working prototype: a real database, real authentication, real security policies, and a live AI call run in the demo. Nothing on screen is hard-coded.

Q: Can't anyone just call ChatGPT?

The value isn't the model, it's the pipeline around it: schema-validated structured output, tolerant parsing, employee-mismatch rejection, transparent activity — and then **auto-execution** of that diagnosis into a governed plan.

Q: How do you stop the AI from hallucinating?

The app validates every response against a strict schema before showing it, rejects anything that doesn't match, never shows cached/fabricated results, and surfaces failures honestly.

Q: What if the agent returns nothing or times out?

180-second cap, one automatic retry, then a clear error — the app never invents a diagnosis.

Q: What does “automatically executes the solution” actually do?

It auto-orchestrates inside the app: creates the case, turns the agent's recommended actions into tasks, assigns the suggested owners, sets staggered due dates, and stores the plan summary. External dispatch is honestly labeled “not connected” — we don't fake it.

Q: Who is responsible for the actions?

The human team. The plan is automatic; completing a task requires recorded evidence by an accountable person. The agent is decision support, not an employment decision-maker.

Q: How is security handled?

Row-Level Security on every table, all mutations through server-side functions, service-role access never in the client. The tests prove direct writes and forged ownership are denied.

Q: What data does the agent see?

Only the selected employee's stored record. If the response doesn't match that employee, it is rejected.

Q: How would you scale this?

Add external integrations as adapters on the same backend functions, swap/tune models through the gateway, and add notifications so plans execute over time, not just at creation.

Q: Why no real HRIS integration?

Honest scope: the execution contract is built and the integrations are adapters. We prioritized a truthful closed loop over a fake integration.

Q: What was the hardest technical problem?

Reliability and honesty of live AI output — envelope variance, malformed JSON, latency. Solved with normalization, JSON repair, balanced extraction, retries — and then closing the loop without fabricating execution.

Q: How is this different from a checklist tool like Trello?

The diagnosis is generated per employee and the plan is derived from it — tasks carry owners and due dates straight from the agent, tied to risk telemetry, with evidence-based tracking and history.

8 Pitfalls to avoid

- **Don't** spend more than ~30 seconds on any non-prototype screen.
- **Don't** claim integrations (HRIS/LMS/email) exist — they are labeled “not connected”.
- **Don't** improvise numbers. “About a minute for the agent” is fine; exact percentages are not.
- **Don't** say “the AI is always right” — say it's decision support with guardrails.
- **Do** stop at 5:00 and own the questions. The questions are where critical thinking scores.

RetainIQ · Final presentation team prep · Based on the working prototype, not slides · Practice the demo twice with the timer before the day · Demoer: open the highest-risk employee first.